

iu-ne-lab-05-Iskander_Nafikov

- **Name:** Iskander Nafikov
- **E-mail:** i.nafikov@innopolis.university
- **Username:** [iskanred](#)
- **Hostname:** lenovo

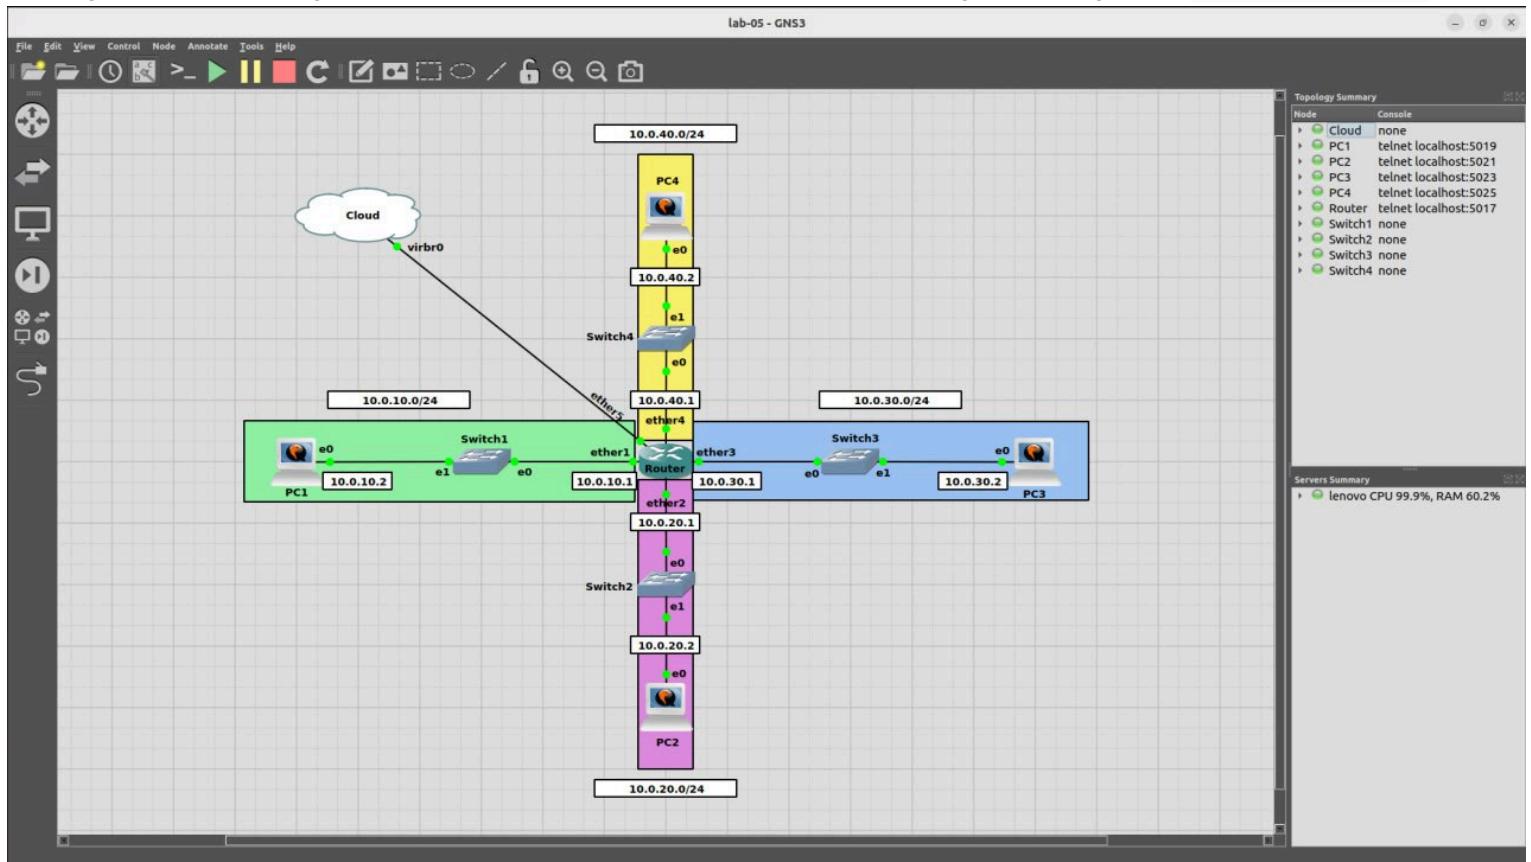
Task 1 - Prepare your network topology

1. In the GNS3 project, select and install a virtual routing solution that you would like to use: Mikrotik (recommended), Pfsense, vyos and so on.'

I decided to use Mikrotik as recommended - Mikrotik CHR 7.16

2. Prepare a simple network consisting of at least one router and two hosts. About four hosts in the network are most optimal. You also might need Internet access for the hosts. It may be something like this, just for the imagination:

I configured the following network. The internet access from PCs was configured using NAT (ip firewall nat).



Task 2 - QoS learning & configuring

1. Let's start with a little theory. Briefly answer the questions or give one-line description what is it: Class of Service (CoS), ToS (Type Of Service), Differentiated Services Code Point (DSCP), Serialization, Packet Marking, Tail Drop, Head Drop, The Leaky bucket algorithm, The Token Bucket Algorithm, Traffic shaping, Traffic policing

- **Class of Service (CoS)**: A mechanism for prioritizing network traffic by classifying packets into different categories for better Quality of Service (QoS), and it is implemented in Link Layer having special 3-bit field in 802.1Q header.
- **Type of Service (ToS)**: An older field in the IP header used to indicate the priority and treatment of packets in terms of delay, throughput, and reliability.
- **Differentiated Services Code Point (DSCP)**: A field in the IP header used for marking packets to indicate the desired level of service, enabling QoS strategies.
- **Serialization**: The process of converting data structures or objects into a format suitable for transmission over a network.
- **Packet Marking**: The practice of labeling packets with specific information (like priority) to influence their treatment as they traverse the network.
- **Tail Drop**: A queuing discipline where packets are dropped from the end of the queue when it reaches its maximum capacity.
- **Head Drop**: A queuing discipline where packets are dropped from the front of the queue, which can be used in certain network management strategies, though it's less common than tail drop.
- **The Leaky Bucket Algorithm**: A traffic shaping method that controls data output rates by allowing data to "leak" out at a constant rate, smoothing bursts.
- **The Token Bucket Algorithm**: A traffic shaping technique that allows for bursts of data while maintaining a defined average rate by using tokens to regulate packet transmission.
- **Traffic Shaping**: The technique of controlling the volume of traffic being sent into a network to optimize or guarantee performance, improve latency, or manage bandwidth.
- **Traffic Policing**: A method of monitoring network traffic and enforcing bandwidth limits by dropping or marking packets that exceed the agreed rate.

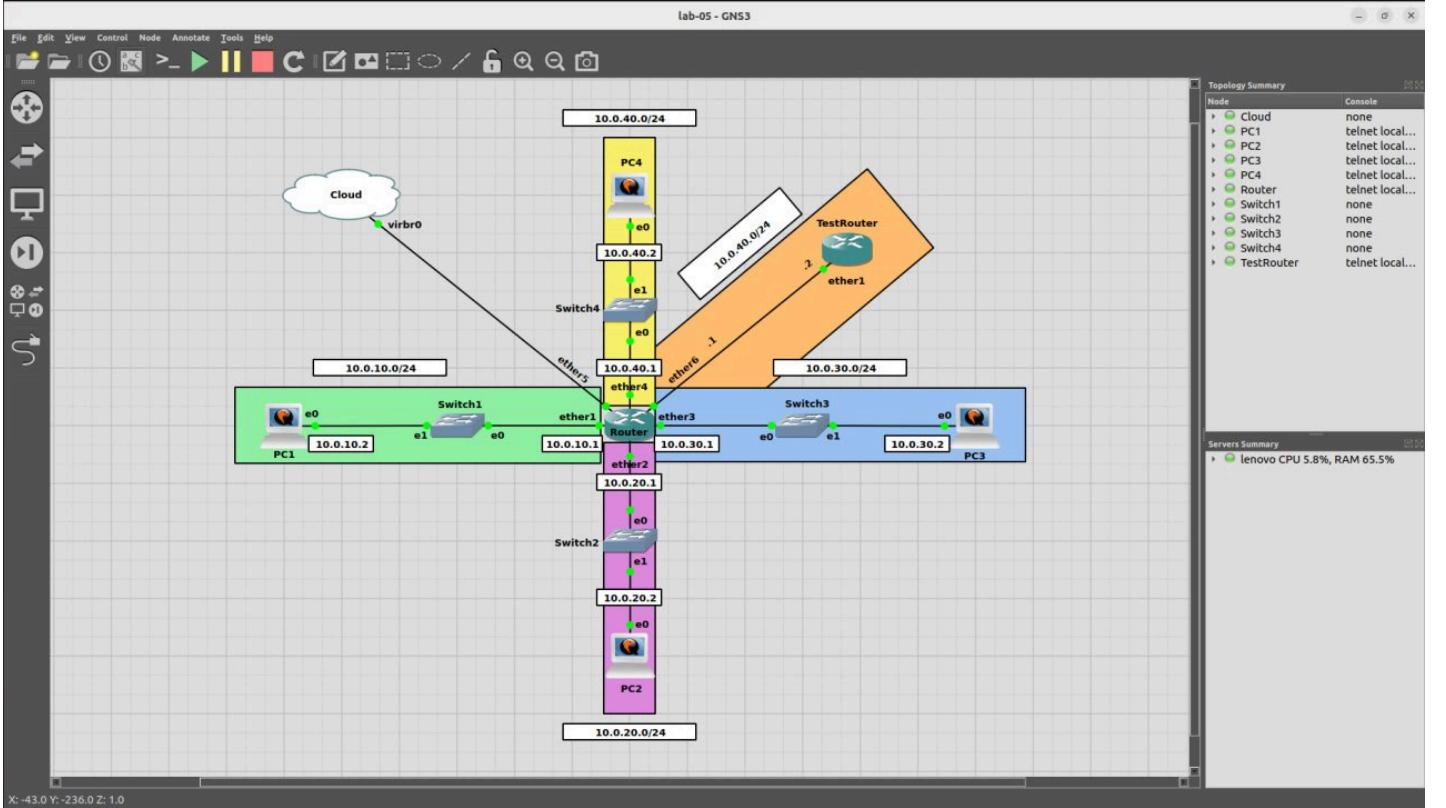
2. Configure your network as you decided above. After your network is configured (don't forget to show the main configuration steps in the report), try to set a speed limitation (traffic shaping) between the two hosts.

I configured my network in the **Task 1**. So now, let define the speed limits.

1. But before, let's investigate what is the current maximum speed without any explicitly set limits.
 - First, it is written in Mikrotik documentation Cloud Host Router (CHR), which is "a RouterOS version intended for running as a virtual machine", has a speed limit 1Mbit for a free version.

CHR Licensing		
The CHR (Cloud Hosted Router) has 4 license levels:		
• free		
• p1 perpetual-1 (\$45)		
• p10 perpetual-10 (\$95)		
• p-unlimited perpetual-unlimited (\$250)		
The 60-day free trial license is available for all paid license levels. To get the free trial license, you have to have an account on MikroTik.com as all license management is done there.		
Perpetual is a lifetime license (buy once, use forever). It is possible to transfer a perpetual license to another CHR instance. A running CHR instance will indicate the time when it has to access the account server to renew its license. If the CHR instance will not be able to renew the license it will behave as if the trial period has run out and will not allow an upgrade of RouterOS to a newer version.		
After licensing a running trial system, you must manually run the /system license renew function from the CHR to make it active. Otherwise, the system will not know you have licensed it in your account. If you do not do this before the system deadline time, the trial will end and you will have to do a complete fresh CHR installation, request a new trial, and then license it with the license you had obtained.		
License	Speed limit	Price
Free	1Mbit	FREE
P1	1Gbit	\$45
P10	10Gbit	\$95
P-Unlimited	Unlimited	\$250

- Let's verify it
- To compute the bandwidth of the router and/or of a link, I added another one router `TestRouter` and connect it to my main one named `Router` in a new LAN `10.0.50.0/24`



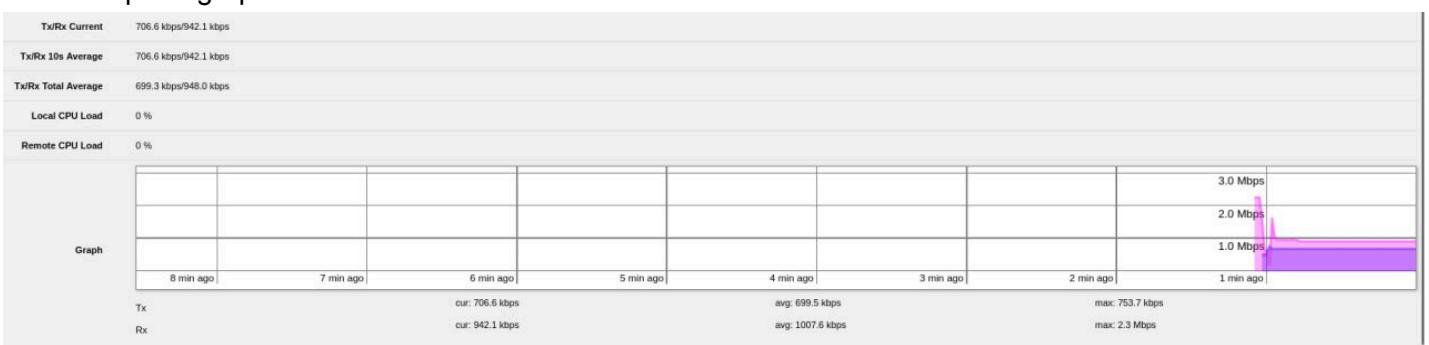
- I did it to run a Mikrotik tool called [bandwidth-tester](#) which requires a server and a client running on Mikrotik device. Let's run speed tests
 - From TestRouter to Router

```
[admin@MikroTik] > tool bandwidth-test direction=both protocol=tcp duration=20s local-tx-speed=50000000 address=10.0.50.1
      status: done testing
      duration: 21s
      tx-current: 938.3kbps
      tx-10-second-average: 937.1kbps
      tx-total-average: 1006.9kbps
      rx-current: 880.3kbps
      rx-10-second-average: 926.7kbps
      rx-total-average: 1013.6kbps
      random-data: no
      direction: both
      connection-count: 20
      local-cpu-load: 2%
      remote-cpu-load: 0%
```

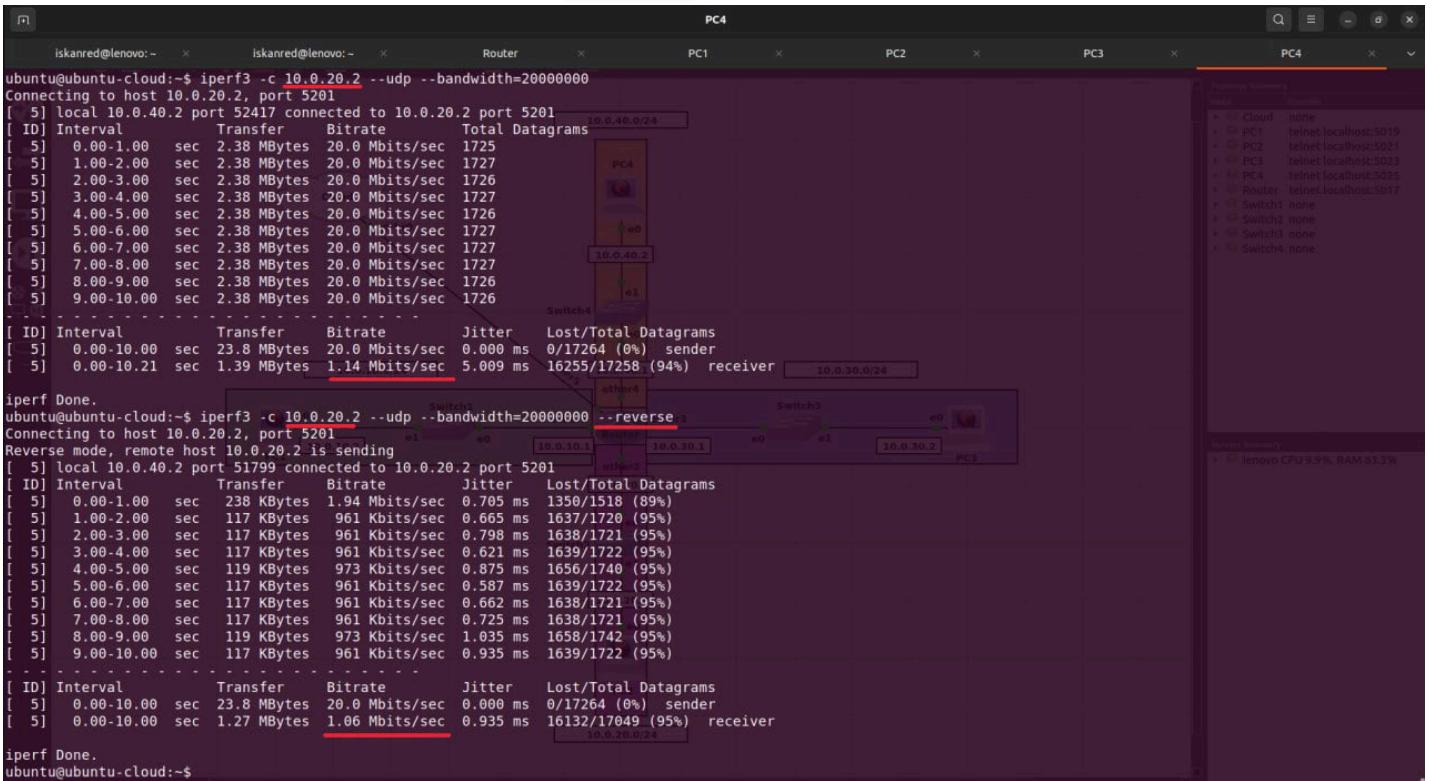
- From Router to TestRouter

```
[admin@MikroTik] > tool bandwidth-test direction=both protocol=tcp duration=20s local-tx-speed=50000000 address=10.0.50.2
      status: done testing
      duration: 21s
      tx-current: 949.9kbps
      tx-10-second-average: 933.6kbps
      tx-total-average: 994.1kbps
      rx-current: 915.1kbps
      rx-10-second-average: 929.0kbps
      rx-total-average: 971.9kbps
      random-data: no
      direction: both
      connection-count: 20
      local-cpu-load: 3%
      remote-cpu-load: 0%
```

- And the speed graph which I took from the Mikrotik Web

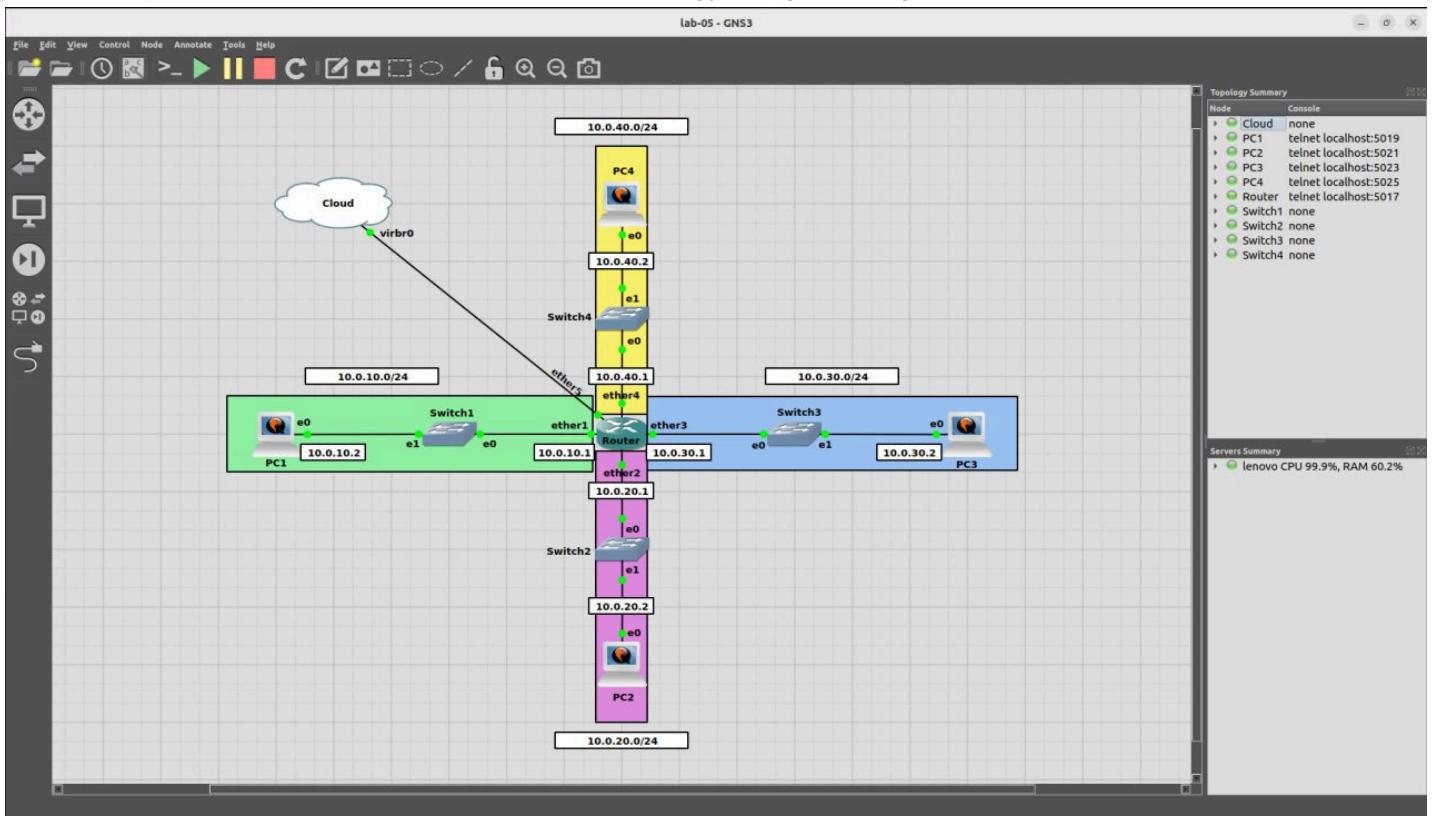


- To verify it completely I run performance test using iperf from PC4 as a client to PC2 as a server at this time using UDP to avoid losing sending or receiving speed due to TCP's congestion control and its bidirectional nature. To make it more convincing I ran performance testing 2 times: from the client (PC4) to the server (PC2) and vice a versa from the server to the client using an option --reverse (--bidir was not an option since it used to me [wrong results](#)). Also, I set --bandwidth equals to 20 Mbit/s to be sure that the generating bitrate won't influence on the result (by default it is 1 Mbit/s which may confuse).



- We see that the bitrate is almost 1 Mbit . Considering computational errors over this time interval it is easy to conclude that 1 Mbit is a hard-limit.
- **Conclusion:** Finally, we can conclude that the **limit actually exists**, it is configured to 1 Mbit and it is configured inside a router's firmware. To increase the limit users should purchase paid versions. Therefore, let's continue using the free version and 1 Mbit limit.

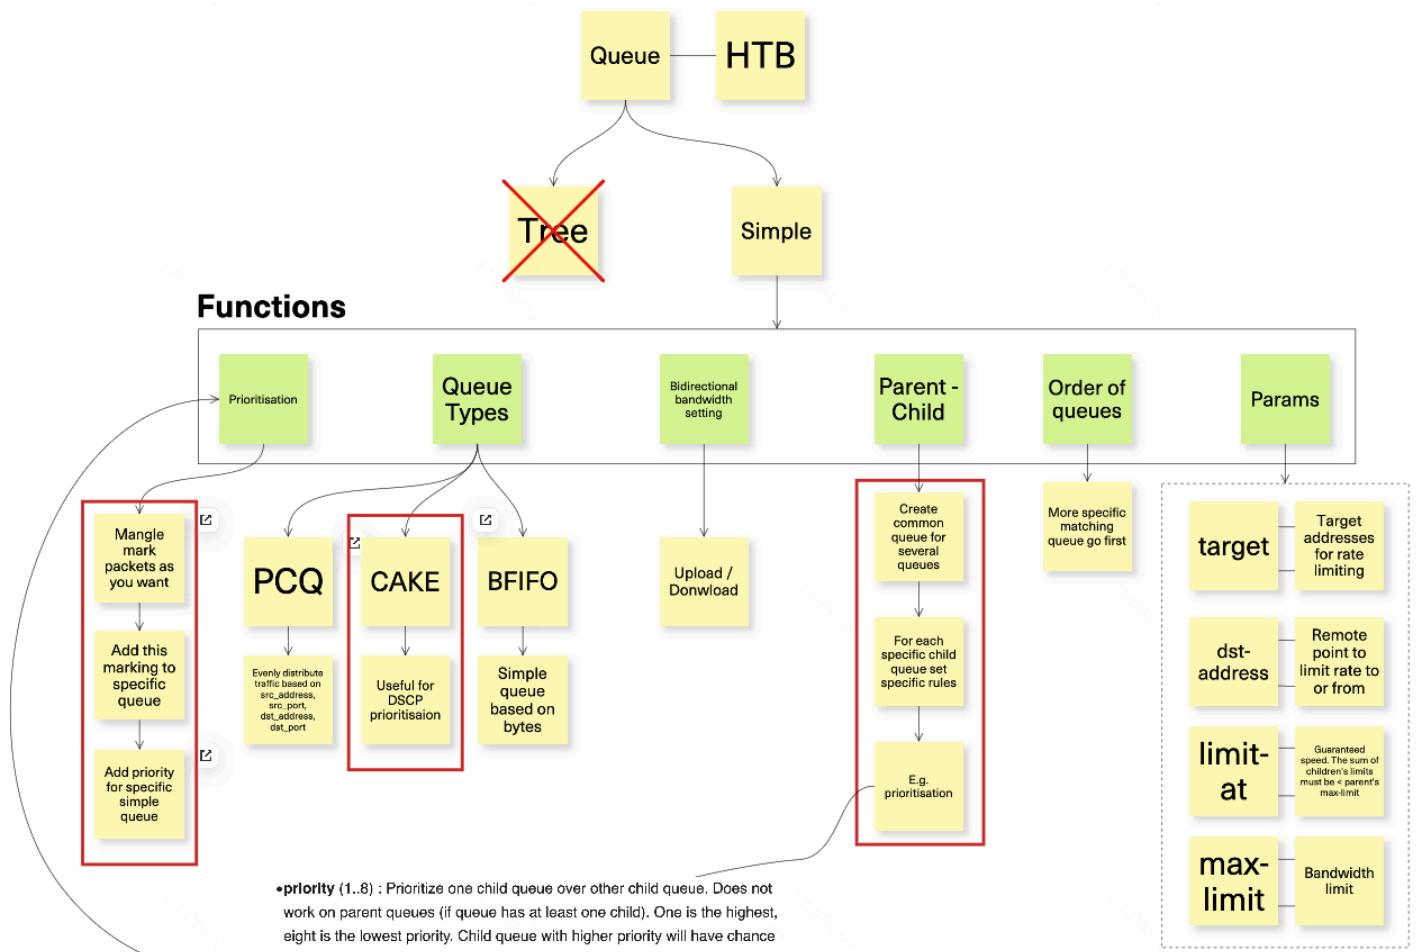
- I removed TestRouter and, thus, returned the topology's original image



2. Now let's set the speed limits using the traffic shaper on our router for the hosts PC3 and PC3

- I decided to apply traffic shaper to my Mikrotik device for connections between PC1 - PC3 .
- To set traffic shaping or scheduling I had to use queues, so I decided to learn something about them and created such a graphical document for myself to structure my findings and keep my knowledge. I decided not to deep dive into queue tree or Hierarchical Token Buckets HTB topics because they are too difficult for me now. Therefore, I tried to keep details about queue simple . This block may not be detailed or totally correct,

but it is clear for me.



- For this purpose I firstly configured a new queue type `full-shaper` based on the `pfifo` (default packet FIFO queue) and set queue size (`pfifo-limit`) to 1 packet making it almost a full 100% shaper. This means that almost every packet that goes over the predefined limit will be dropped immediately. Therefore, it will cause a **big packet loss**, but **low-latency packets delivery**. I decided to use `pfifo` just because of its simplicity.

```
/queue type add name="full-shaper" kind=pfifo pfifo-limit=1
```

```
10 * name="default-small" kind=pfifo pfifo-limit=10
[admin@MikroTik] > queue typ print where name=full-shaper
Flags: * - default
5   name="full-shaper" kind=pfifo pfifo-limit=1
[admin@MikroTik] >
```

- Then I created a simple queue. I decided to use queue simple just for the sake of its simplicity. I configured a queue in the following format:

```
/queue simple add name="host_1-3_parent_queue" target="10.0.30.2" dst="10.0.10.2"
queue="full_shaper/full_shaper" max-limit="200K/400K"
```

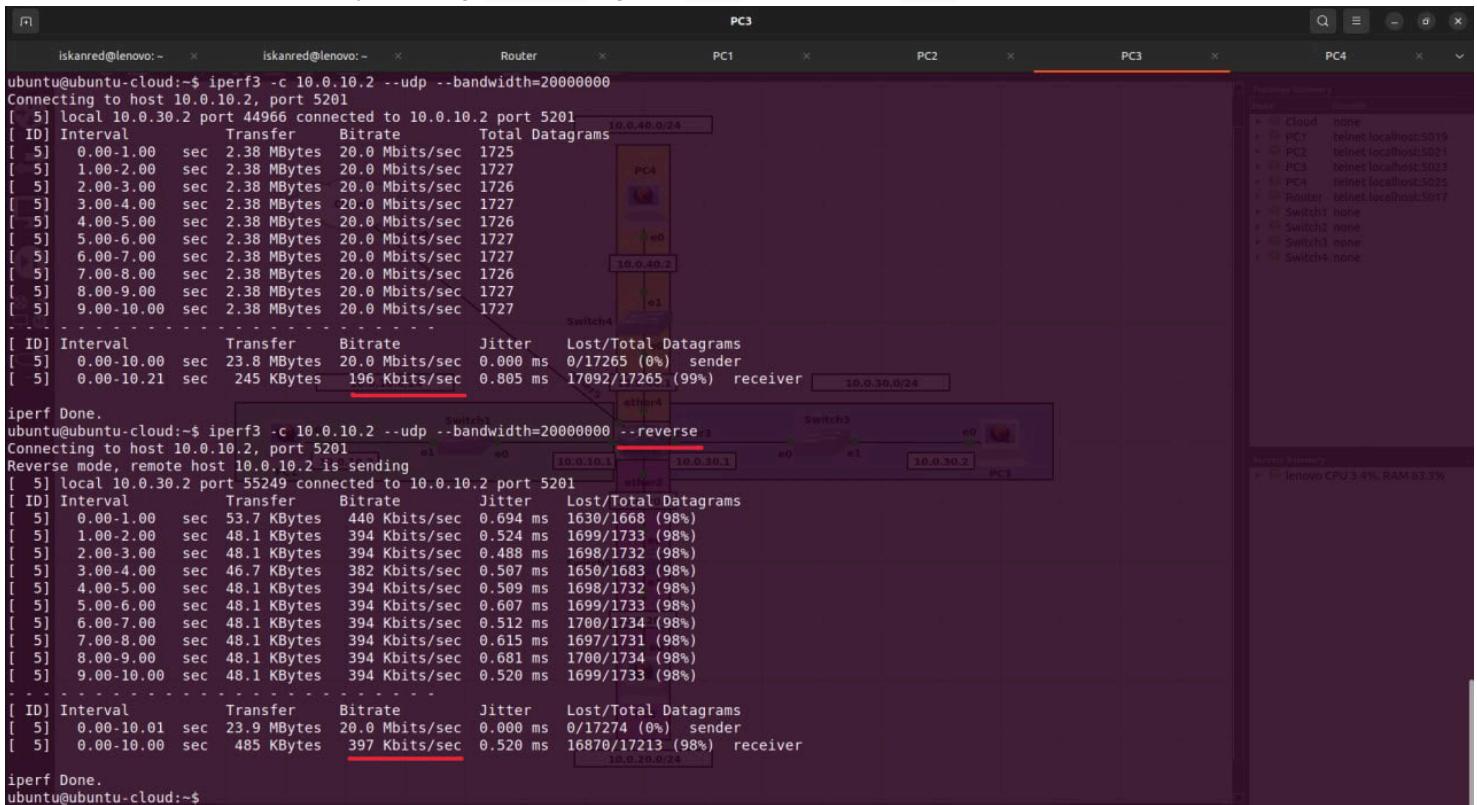
```
[admin@MikroTik] > queue simple print
Flags: X - disabled, I - invalid; D - dynamic
0   name="host_1-3_parent_queue" target=10.0.30.2/32 dst=10.0.10.2/32 parent=none packet-marks="" priority=8/8
queue=full-shaper/full-shaper limit-at=0/0 max-limit=200k/400k burst-limit=0/0 burst-threshold=0/0 burst-time=0s/0s
bucket-size=0.1/0.1
[admin@MikroTik] >
```

- As you can I configured a queue `host_1-3_parent_queue` for the link between PC3 (`target=10.0.30.2`) and PC1 (`dst=10.0.10.2`) with the queue type `queue=full-shaper` for both downloading and uploading

traffic and the actual max bandwidth of 200 Kbit/s for uploading and of 400 Kbit/s for downloading (from the client's perspective).

- 3. Run a bandwidth testing tool, see what is the max speed you can get and verify your speed limitation. Compare the speed between the different hosts.

- Now let's check if it works by running `iperf3` again (PC1 is a server, PC3 is a client, UDP traffic)



```

ubuntu@ubuntu-cloud:~$ iperf3 -c 10.0.10.2 --udp --bandwidth=20000000
Connecting to host 10.0.10.2, port 5201
[ 5] local 10.0.30.2 port 44966 connected to 10.0.10.2 port 5201
[ ID] Interval Transfer Bitrate Total Datagrams
[ 5] 0.00-1.00 sec 2.38 MBytes 20.0 Mbits/sec 1725
[ 5] 1.00-2.00 sec 2.38 MBytes 20.0 Mbits/sec 1727
[ 5] 2.00-3.00 sec 2.38 MBytes 20.0 Mbits/sec 1726
[ 5] 3.00-4.00 sec 2.38 MBytes 20.0 Mbits/sec 1727
[ 5] 4.00-5.00 sec 2.38 MBytes 20.0 Mbits/sec 1726
[ 5] 5.00-6.00 sec 2.38 MBytes 20.0 Mbits/sec 1727
[ 5] 6.00-7.00 sec 2.38 MBytes 20.0 Mbits/sec 1727
[ 5] 7.00-8.00 sec 2.38 MBytes 20.0 Mbits/sec 1726
[ 5] 8.00-9.00 sec 2.38 MBytes 20.0 Mbits/sec 1727
[ 5] 9.00-10.00 sec 2.38 MBytes 20.0 Mbits/sec 1727
[ 5] 0.00-10.00 sec 23.8 MBytes 20.0 Mbits/sec 1727
[ 5] 0.00-10.21 sec 245 KBytes 196 Kbits/sec 0.805 ms 17092/17265 (99%) receiver 10.0.30.0/24

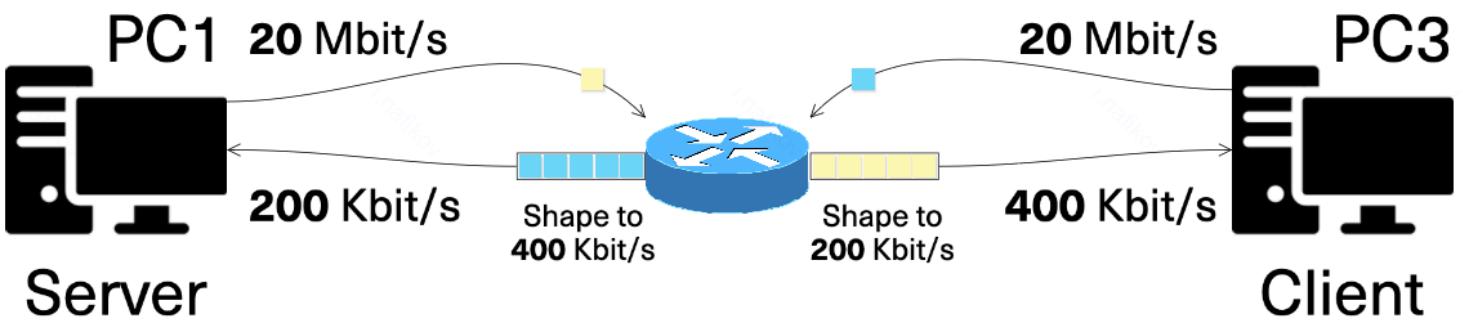
[ ID] Interval Transfer Bitrate Jitter Lost/Total Datagrams
[ 5] 0.00-10.00 sec 23.8 MBytes 20.0 Mbits/sec 0.000 ms 0/17265 (0%) sender
[ 5] 0.00-10.21 sec 245 KBytes 196 Kbits/sec 0.805 ms 17092/17265 (99%) receiver 10.0.30.0/24

iperf Done.
ubuntu@ubuntu-cloud:~$ iperf3 -c 10.0.10.2 --udp --bandwidth=20000000 --reverse
Connecting to host 10.0.10.2, port 5201
Reverse mode, remote host 10.0.10.2 is sending
[ 5] local 10.0.30.2 port 55249 connected to 10.0.10.2 port 5201
[ ID] Interval Transfer Bitrate Jitter Lost/Total Datagrams
[ 5] 0.00-1.00 sec 53.7 KBytes 440 Kbits/sec 0.694 ms 1630/1668 (98%)
[ 5] 1.00-2.00 sec 48.1 KBytes 394 Kbits/sec 0.524 ms 1699/1733 (98%)
[ 5] 2.00-3.00 sec 48.1 KBytes 394 Kbits/sec 0.488 ms 1698/1732 (98%)
[ 5] 3.00-4.00 sec 46.7 KBytes 382 Kbits/sec 0.507 ms 1650/1683 (98%)
[ 5] 4.00-5.00 sec 48.1 KBytes 394 Kbits/sec 0.509 ms 1698/1732 (98%)
[ 5] 5.00-6.00 sec 48.1 KBytes 394 Kbits/sec 0.607 ms 1699/1733 (98%)
[ 5] 6.00-7.00 sec 48.1 KBytes 394 Kbits/sec 0.512 ms 1700/1734 (98%)
[ 5] 7.00-8.00 sec 48.1 KBytes 394 Kbits/sec 0.615 ms 1697/1731 (98%)
[ 5] 8.00-9.00 sec 48.1 KBytes 394 Kbits/sec 0.681 ms 1700/1734 (98%)
[ 5] 9.00-10.00 sec 48.1 KBytes 394 Kbits/sec 0.520 ms 1699/1733 (98%)

[ ID] Interval Transfer Bitrate Jitter Lost/Total Datagrams
[ 5] 0.00-10.01 sec 23.9 MBytes 20.0 Mbits/sec 0.000 ms 0/17274 (0%) sender
[ 5] 0.00-10.00 sec 485 KBytes 397 Kbits/sec 0.520 ms 16870/17213 (98%) receiver 10.0.20.0/24

iperf Done.
ubuntu@ubuntu-cloud:~$
```

- We see that it worked and the actual bitrate was definitely about 200 Kbit/s and 400Kbit/s for uploading and downloading correspondingly. The scheme of interaction look like:



- Nevertheless, we can notice that the packet loss is enormous: 99% for the client's packets and 98% for the server's packets.

```

Switch4
Lost/Total Datagrams
0/17265 (0%) sender
17092/17265 (99%) receiver
eth0
200000000 --reverse
10.0.10.1 10.0.30.1
5201
Lost/Total Datagrams
1630/1668 (98%)
1699/1733 (98%)
1698/1732 (98%)
1650/1683 (98%)
1698/1732 (98%)
1699/1733 (98%)
1700/1734 (98%)
1697/1731 (98%)
1700/1734 (98%)
1699/1733 (98%)
Lost/Total Datagrams
0/17274 (0%) sender
16870/17213 (98%) receiver
10.0.20.0/24

```

- Such number of lost datagrams happened because UDP is not reliable, especially when the bandwidth and the actual bitrate differ so much. Also, such big numbers occurred because of very aggressive traffic shaping which drops packets almost immediately with queue size = 1 packet.
- Therefore, we can tune it a bit and really slightly decrease packet loss configuring our shaper to be more a scheduler with the queue size = 300 packets.

```
/queue type set full-shaper pfifo-limit=300
```



The screenshot shows the Winbox interface with several windows open:

- Router**: Shows the network topology with nodes: cloud, PC1, PC2, PC3, PC4, Router, Switch1, and Switch2.
- PC1**, **PC2**, **PC3**, and **PC4**: Individual terminal windows for each host.
- Router**: A terminal window showing the configuration of the Router's queueing discipline.
- Terminal**: A terminal window showing the command history.

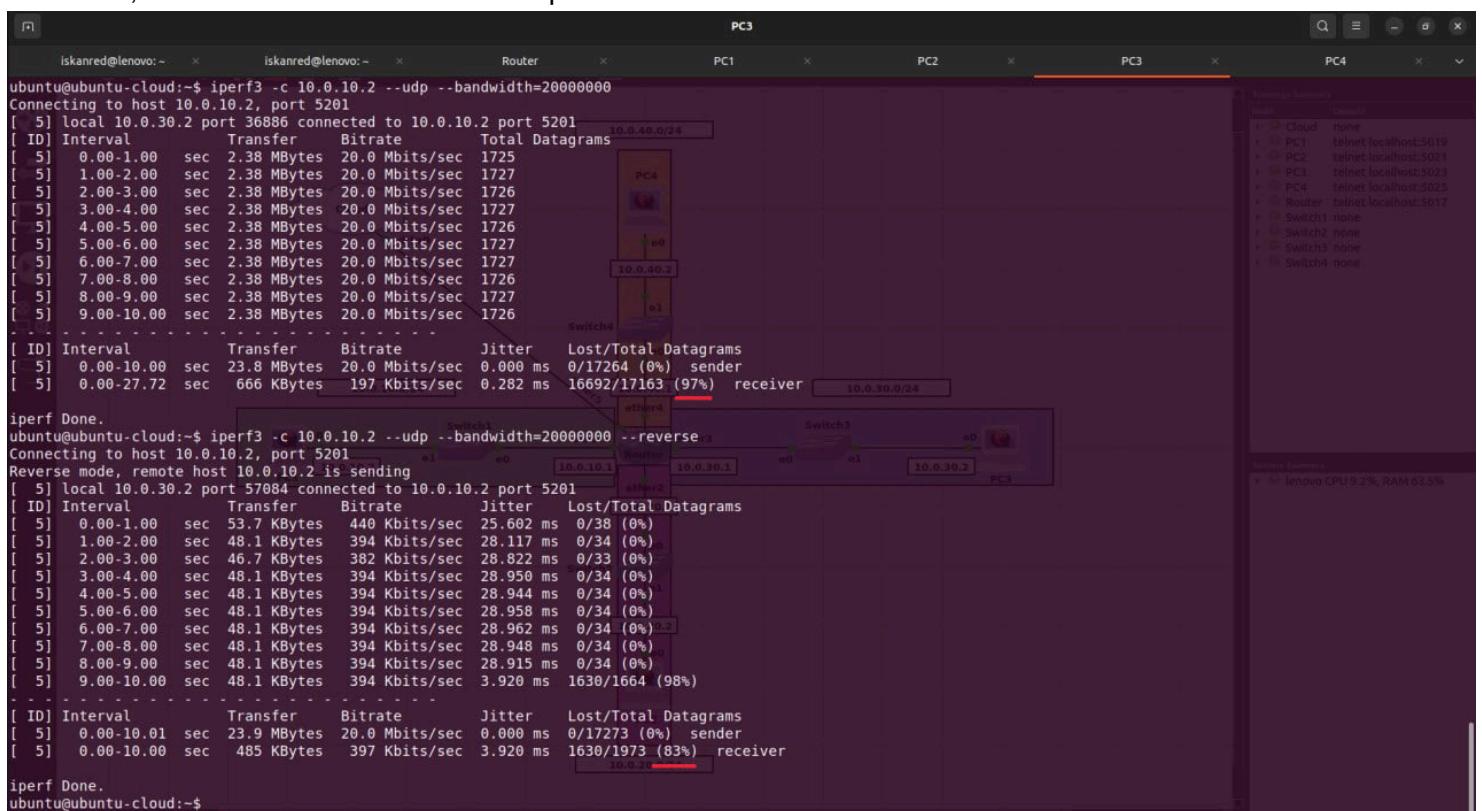
In the Router window, the configuration command is:

```
[admin@MikroTik] > queue type set full-shaper pfifo-limit=300
```

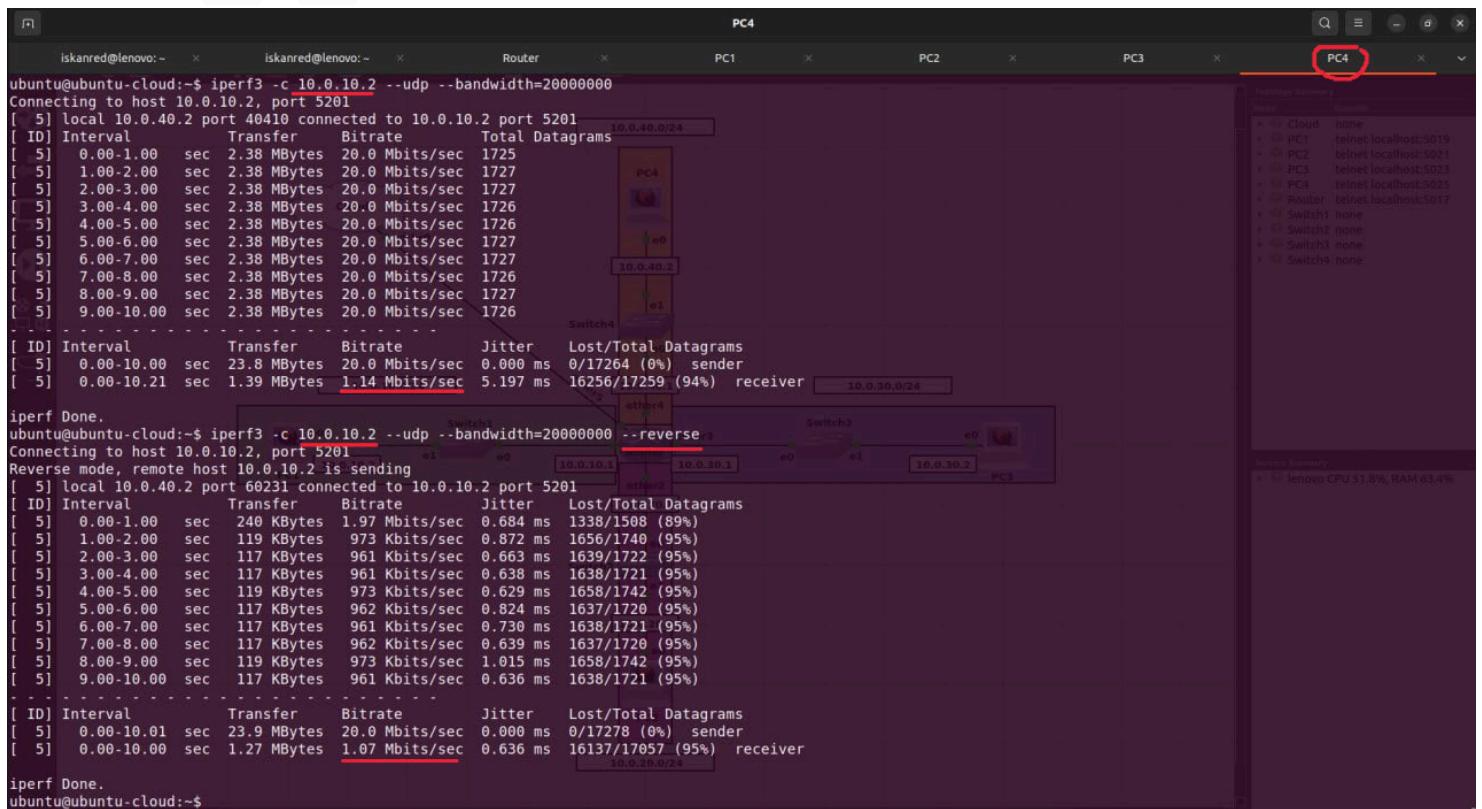
This command changes the queue discipline for the "full-shaper" from a queue size of 1 to 300. The configuration output shows:

```
Flags: * - default
5 | name="full-shaper" kind=pfifo pfifo-limit=1
[admin@MikroTik] > queue type set full-shaper pfifo-limit=300
[admin@MikroTik] > queue type print where name=full-shaper
Flags: * - default
5 | name="full-shaper" kind=pfifo pfifo-limit=300
```

- Afterward, we can notice this decrease on practice: 99% -> 97% 98% -> 83%



- However, the delay was increased also. I had to wait longer for the end of test.
- Still the speed for other clients remained the same = 1 Mbit/s even for the same server PC1 . For example for the connection PC4 -- PC1 :



4. While your bandwidth test is still running, try to download a file from one host to the other host and see what is the max speed you can get. If you have more than two hosts on the network, play around with different speed values and show it.

- To complete this task I set up an FTP server on PC1 . I used `pure-ftpd` for FTP server

```
File Edit View Search Terminal Tabs Help
iskanred@lenovo:~> iskanred@lenovo:~> Router PC1 PC2 PC3 PC4
ubuntu@ubuntu-cloud:~$ nmap -p - localhost
Starting Nmap 7.94SVN ( https://nmap.org ) at 2024-12-15 19:10 UTC
Nmap scan report for localhost (127.0.0.1)
Host is up (0.00012s latency).
Not shown: 65533 closed tcp ports (conn-refused)
PORT      STATE SERVICE
21/tcp    open  ftp
22/tcp    open  ssh

Nmap done: 1 IP address (1 host up) scanned in 2.19 seconds
ubuntu@ubuntu-cloud:~$
```



The network diagram illustrates a topology with a central 'Cloud' node connected to five other nodes: 'PC1', 'PC2', 'PC3', 'PC4', and 'Router'. Additionally, there are four 'Switch' nodes labeled 'Switch1', 'Switch2', 'Switch3', and 'Switch4', which are also interconnected. The 'PC1' node is highlighted in red, indicating it is the target of the Nmap scan.

- Also I created a file of 50MB on PC1 using this Python script:

```
with open('my_file', 'wb') as f:  
    num_chars = 1024 * 1024 * 1024  
    f.write('0' * num_chars)
```

- Before all, I decided to update my queue and set the limit to 1 Mbit/s for both download and upload. Also, I changed its type to pcq-upload-default/pcq-download-default just because I wanted it to try, but there should be no effect since pcq-classifier is dst-address and it will be the same for both the connections: file download and perf.

```
[admin@MikroTik] > queue simple print
Flags: X - disabled, I - invalid, D - dynamic
0  name="host_1-3.parent_queue" target=10.0.30.2/32 dst=10.0.10.2/32 parent=None
  packet-marks="" priority=8/8 queue=pcq-upload-default/pcq-download-default
  limit-at=0/0 max-limit=1M/1M burst-limit=0/0 burst-threshold=0/0
  burst-time=0s/0s bucket-size=0.1/0.1
[admin@MikroTik] >
```

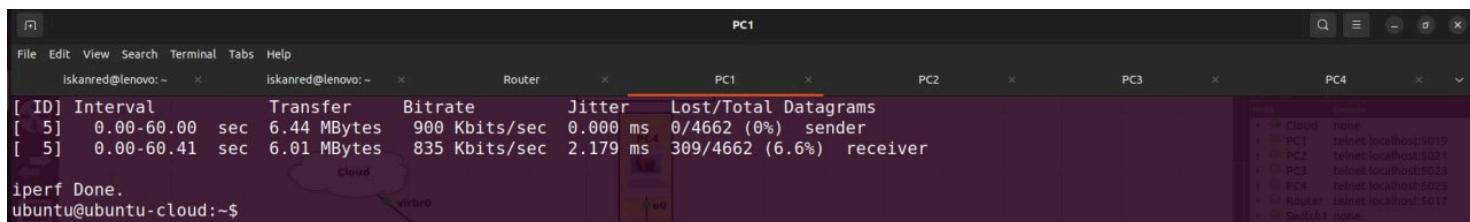
- Then I decided to collect the speed statistics without running perf firstly

```
File Edit View Search Terminal Tabs Help  
ikanred@lenovo:~> iskanred@lenovo:~> Router > PC1 > PC2 > PC3 > PC4  
ubuntu@ubuntu-cloud:~$ curl -O ftp://ubuntu:ubuntu@10.0.10.2/my_file  
% Total % Received % Xferd Average Speed Time Time Current  
Dload Upload Total Spent Left Speed  
0 50.0M 0 162k 0 0 113k 0 0:07:32 0:00:01 0:07:31 113k  
2 50.0M 2 1221k 0 0 117k 0 0:07:16 0:00:10 0:07:06 118k  
4 50.0M 4 2402k 0 0 117k 0 0:07:14 0:00:20 0:06:54 118k  
7 50.0M 7 3584k 0 0 117k 0 0:07:14 0:00:30 0:06:44 124k  
7 50.0M 7 4013k 0 0 117k 0 0:07:14 0:00:34 0:06:40 118k ^C
```

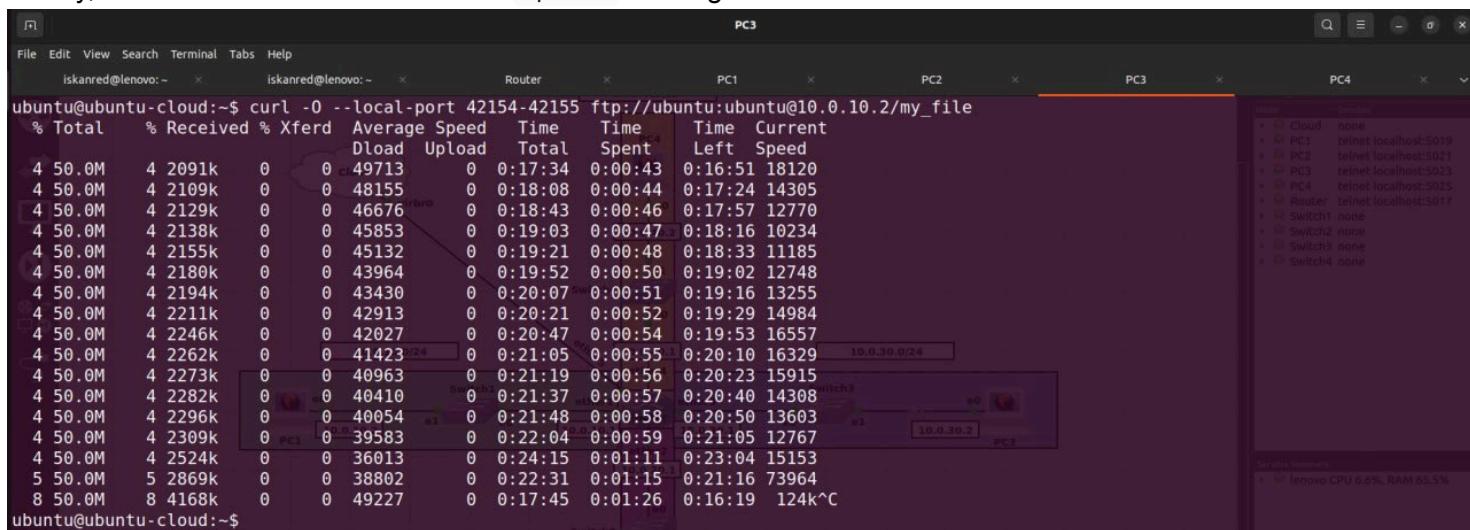
- We see that normally even with 1 Mbit/s open channel the speed is only about 120 Kbit/s. This can happen due to TCP traffic has different traffic guarantees which impact the final speed.

- Then I ran iperf3 from PC1 to PC3 : UDP traffic with regular bandwidth = 900 Kbit/s , not 1 Mbit/s to leave some "space" for FTP traffic.

```
iperf3 -c 10.0.10.2 --bandwidth=300000 --udp --time=60
```



- Finally, I started FTP file download with iperf3 running

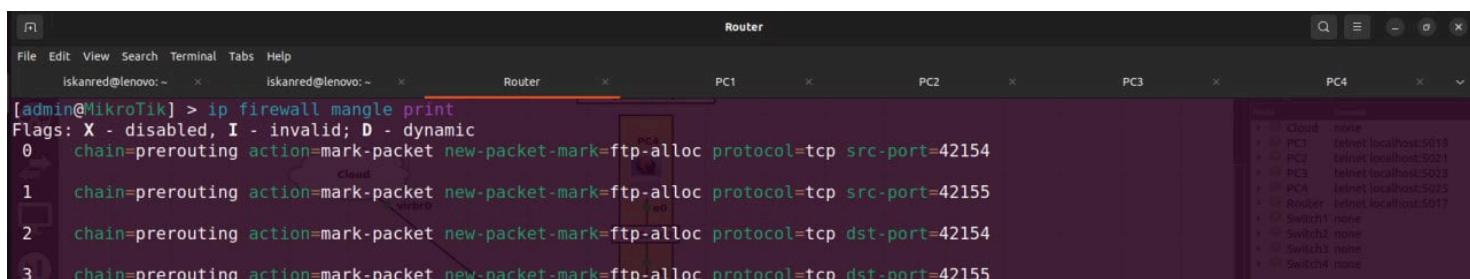


- We can notice that the speed **dropped dramatically** to approximately 15 Kbit/s and returned to 120 Kbit/s only after perf ended.

5. Deploy and verify your QoS rules to prioritise the downloading of a file (or any other scenario) over the bandwidth test.

- Now let me show my QoS configuration. I use traffic allocation technique: limit-at , child queue, packet marks and etc. Let me show it now
- First, I configured my ip firewall mangle to mark all FTP-related packets as ftp-alloc

```
ip firewall mangle>
add chain=prerouting action=mark-packet new-packet-mark=ftp-alloc protocol=tcp src-port=42154
add chain=prerouting action=mark-packet new-packet-mark=ftp-alloc protocol=tcp src-port=42155
add chain=prerouting action=mark-packet new-packet-mark=ftp-alloc protocol=tcp dst-port=42154
add chain=prerouting action=mark-packet new-packet-mark=ftp-alloc protocol=tcp dst-port=42155
```

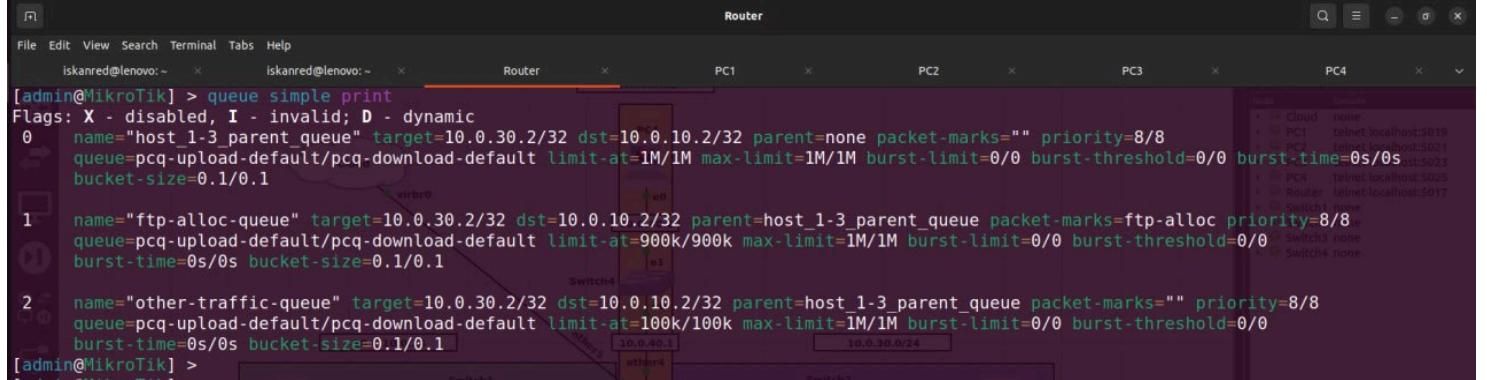


- As you can see I added packet marking on prerouting chain (Rules in this chain apply to packets as they just arrive on the network interface) on TCP protocol for different sets of dst-port and src-port . But why such

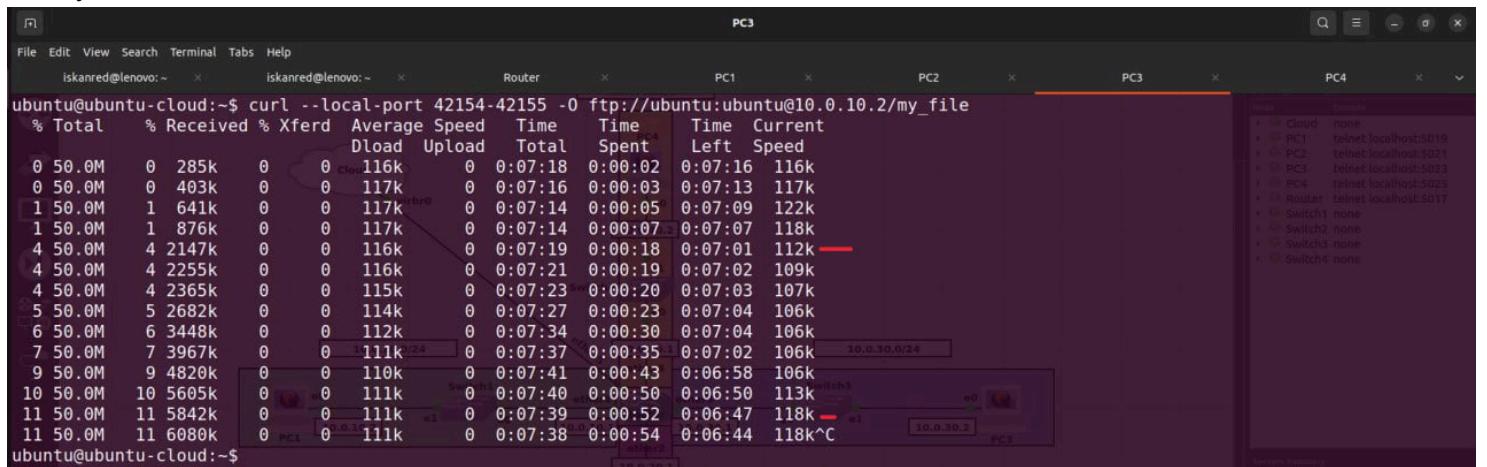
strange numbers: 42154, 42155? This is because I decided to run my FTP client on these ports (one port for a control channel and one for a data channel).

- Then, I added new queues that I made children of my main host_1-3_parent_queue .

```
queue simple>
add name=ftp-alloc-queue target=10.0.30.2 dst=10.0.10.2 parent=host_1-3_parent_queue packet-
marks=ftp-alloc queue=pcq-upload-default/pcq-download-default limit-at=900k/900k max-
limit=1M/1M
add name=ftp-alloc-queue target=10.0.30.2 dst=10.0.10.2 parent=host_1-3_parent_queue packet-
marks=ftp-alloc queue=pcq-upload-default/pcq-download-default limit-at=900k/900k max-
limit=1M/1M
```

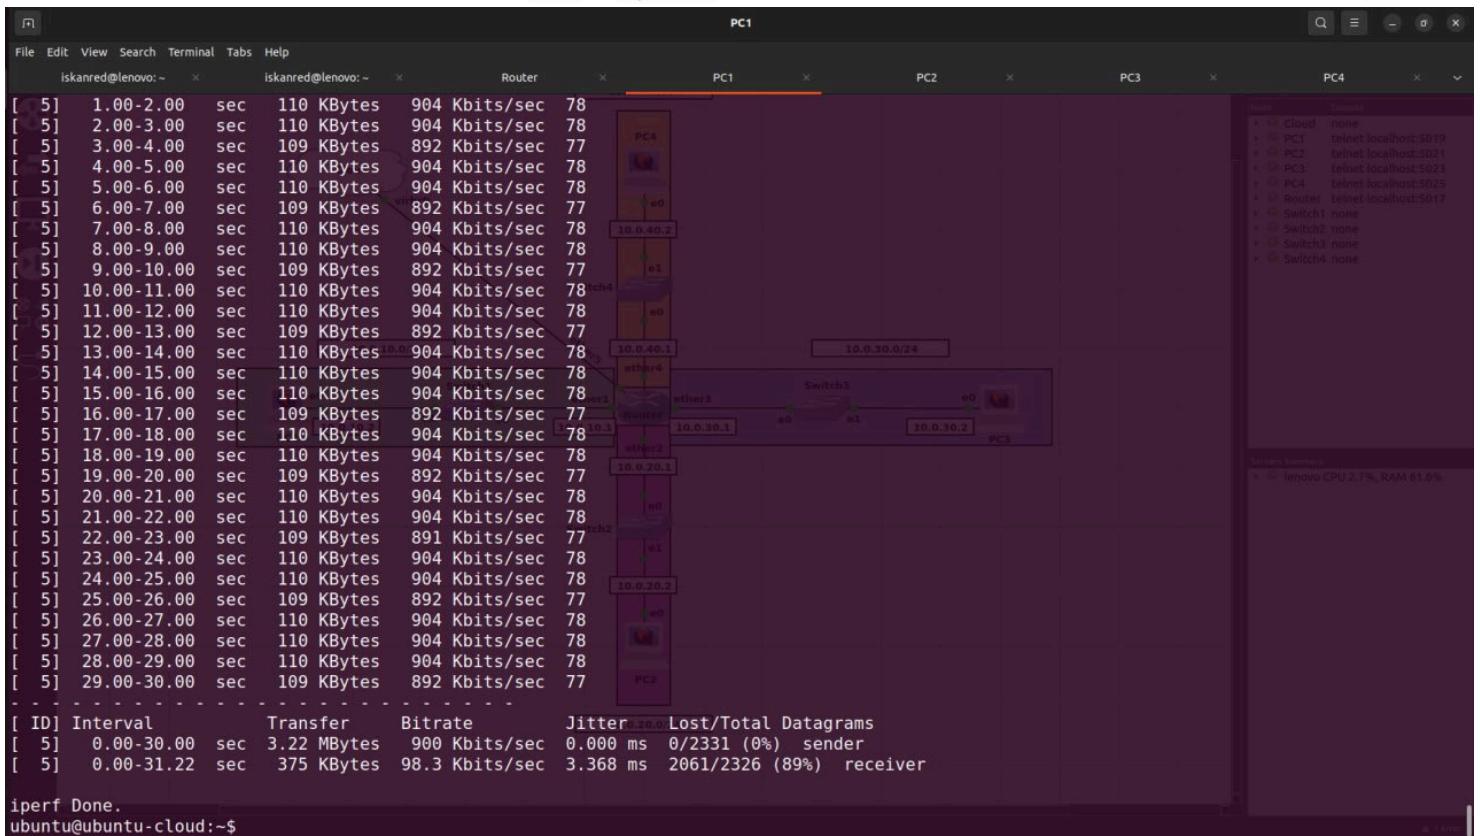


- The first queue ftp-alloc-queue is configured to be a queue for all the packets that are marked as ftp-alloc => all FTP-related packets between these hosts. This queue has limit-at=900k/900k meaning that despite all other traffic flows in other queues, this queue will get guaranteed 900 Kbit/s for both upload/download while the maximum limit leaves the same 1 Mbit/s .
- The second queue other-traffic-queue is also configured for the traffic between these hosts. However, here limit-at parameter equals to 100k/100k meaning that this queue would get only guaranteed 100 Kbit/s for both upload/download. The traffic would go to this queue only if previous queues were not convenient. Therefore, all non-FTP traffic would go to this queue.
- It's worth to say that in this case it's unnecessary to have a parent queue. However, it's more flexible.
- Finally, we can run our FTP file download



- It's easy to notice that the speed drop in this case is really small, about 12 Kbit/s from 118 Kbit/s to 106 Kbit/s

- It worked! And we can double check it in PC1's iperf3 stats



- We see that the receiver's bitrate did not exceed 100 Kbit/s

6. What is the difference between the QoS rules to traffic allocation and priority-based QoS? Try to set up each of them and show them. In which tasks of this lab do you use one or the other?

- **Traffic allocation** is a technique to define specific policies to control how much bandwidth is assigned to different types of traffic or users within a network. The example is above for downloading files with FTP. It is an efficient way to statically give specific available bandwidth!
- **Priorities** are used classifying packets and assigning them different priorities to ensure critical applications receive better bandwidth and low-latency treatment, regardless of overall usage. Here we cannot guarantee some specific number as a bandwidth value. However, it is better for low-latency applications (such as real-time) since marked packets have a higher probability to be passed first.

Priority-based QoS configuration:

- Since I already deployed traffic allocation technique in the previous task, I decided to deploy priorities using the same queues on Mikrotik.
- First, I removed `limit-at` for the `frp-alloc-queue` to make traffic allocation QoS not having an impact on the results. Nevertheless, I gave `limit-at=10k/10k` for the `other-traffic-queue` to allow non-FTP traffic to have at least some guaranteed speed.
- Then I gave the highest priority to the `ftp-alloc-queue` (1 from 8), while the priority for the `other-traffic-queue` remained the same (8 by default).

```

[admin@MikroTik] > queue simple print
Flags: X - disabled, I - invalid; D - dynamic
0 name="host_1-3_parent_queue" target=10.0.30.2/32 dst=10.0.10.2/32 parent=none packet-marks="" priority=8/8
queue=pcq-upload-default/pcq-download-default limit-at=1M/1M max-limit=1M/1M burst-limit=0/0 burst-threshold=0/0 burst-time=0s/0s
bucket-size=0.1/0.1

1 name="ftp-alloc-queue" target=10.0.30.2/32 dst=10.0.10.2/32 parent=host_1-3_parent_queue packet-marks=ftp-alloc priority=1/1
queue=pcq-upload-default/pcq-download-default limit-at=0/0 max-limit=1M/1M burst-limit=0/0 burst-threshold=0/0 burst-time=0s/0s
bucket-size=0.1/0.1

2 name="other-traffic-queue" target=10.0.30.2/32 dst=10.0.10.2/32 parent=host_1-3_parent_queue packet-marks="" priority=8/8
queue=pcq-upload-default/pcq-download-default limit-at=10k/10k max-limit=1M/1M burst-limit=0/0 burst-threshold=0/0 burst-time=0s/0s
bucket-size=0.1/0.1
[admin@MikroTik] >

```

- I ran started FTP file download and noticed that the speed didn't decrease at all!

```

ubuntu@ubuntu-cloud:~$ curl --local-port 42154-42155 -0 ftp://ubuntu:ubuntu@10.0.10.2/my_file
% Total    % Received % Xferd  Average Speed   Time     Time      Current
          Dload  Upload   Total Spent    Spent  Left Speed
0 50.0M    0 448k    0    0  117k    0:07:16  0:00:03  0:07:13  117k
1 50.0M    1 803k    0    0  117k    0:07:15  0:00:06  0:07:09  118k
4 50.0M    4 2454k    0    0  117k    0:07:14  0:00:20  0:06:54  117k
5 50.0M    5 2811k    0    0  117k    0:07:14  0:00:23  0:06:51  118k
8 50.0M    8 4459k    0    0  117k    0:07:14  0:00:37  0:06:37  117k
11 50.0M   11 5753k    0    0  117k    0:07:14  0:00:48  0:06:26  118k
36 50.0M   36 18.1M   0    0  118k    0:07:13  0:02:37  0:04:36  117k^C
ubuntu@ubuntu-cloud:~$

```

- Afterwards, I checked iperf3 stats and found out that it even went wrong at the end because control socket has closed unexpectedly because of too many packets lost.

```

ubuntu@ubuntu-cloud:~$ iperf3 -c 10.0.30.2 --udp --time=30 --bandwidth=900000
Connecting to host 10.0.30.2, port 5201
[ 5] local 10.0.10.2 port 33296 connected to 10.0.30.2 port 5201
[ ID] Interval Transfer Bitrate Total Datagrams
[ 5] 0.00-1.00  sec 110 KBytes 903 Kbytes/sec 78
[ 5] 1.00-2.00  sec 110 KBytes 904 Kbits/sec 78
[ 5] 2.00-3.00  sec 110 KBytes 904 Kbits/sec 78
[ 5] 3.00-4.00  sec 109 KBytes 892 Kbits/sec 77
[ 5] 4.00-5.00  sec 110 KBytes 904 Kbits/sec 78
[ 5] 5.00-6.00  sec 110 KBytes 904 Kbits/sec 78
[ 5] 6.00-7.00  sec 109 KBytes 892 Kbits/sec 77
[ 5] 7.00-8.00  sec 110 KBytes 904 Kbits/sec 78
[ 5] 8.00-9.00  sec 110 KBytes 904 Kbits/sec 78
[ 5] 9.00-10.00 sec 109 KBytes 892 Kbits/sec 77
[ 5] 10.00-11.00 sec 110 KBytes 904 Kbits/sec 78
[ 5] 11.00-12.00 sec 110 KBytes 904 Kbits/sec 78
[ 5] 12.00-13.00 sec 109 KBytes 892 Kbits/sec 77
[ 5] 13.00-14.00 sec 110 KBytes 904 Kbits/sec 78
[ 5] 14.00-15.00 sec 110 KBytes 904 Kbits/sec 78
[ 5] 15.00-16.00 sec 110 KBytes 904 Kbits/sec 78
[ 5] 16.00-17.00 sec 109 KBytes 892 Kbits/sec 77
[ 5] 17.00-18.00 sec 110 KBytes 904 Kbits/sec 78
[ 5] 18.00-19.00 sec 110 KBytes 904 Kbits/sec 78
[ 5] 19.00-20.00 sec 109 KBytes 892 Kbits/sec 77
[ 5] 20.00-21.00 sec 110 KBytes 904 Kbits/sec 78
[ 5] 21.00-22.00 sec 110 KBytes 904 Kbits/sec 78
[ 5] 22.00-23.00 sec 109 KBytes 892 Kbits/sec 77
[ 5] 23.00-24.00 sec 110 KBytes 904 Kbits/sec 78
[ 5] 24.00-25.00 sec 110 KBytes 904 Kbits/sec 78
[ 5] 25.00-26.00 sec 109 KBytes 892 Kbits/sec 77
[ 5] 26.00-27.00 sec 110 KBytes 904 Kbits/sec 78
[ 5] 27.00-28.00 sec 110 KBytes 904 Kbits/sec 78
[ 5] 28.00-29.00 sec 110 KBytes 904 Kbits/sec 78
[ 5] 29.00-30.00 sec 109 KBytes 892 Kbits/sec 77
[ ID] Interval Transfer Bitrate Jitter Lost/Total Datagrams
[ 5] 0.00-30.00 sec 3.22 MBytes 900 Kbits/sec 0.000 ms 0/2331 (0%) sender
[ 5] 0.00-30.00 sec 0.00 Bytes 0.00 bits/sec 0.000 ms 0/0 (0%) receiver
iperf3: error - control socket has closed unexpectedly
ubuntu@ubuntu-cloud:~$

```

- Therefore, we can conclude that it worked perfectly! In fact, without specified limit-at for other-traffic-queue iperf3 couldn't even connect to the PC3
- I checked that even with priority = 7 ftp-traffic-queue already takes all the traffic, so it is necessary to set the limit-at .

DSCP and CAKE

- What is more, I tried to use DSCP with the help of a CAKE queueing discipline.
- Firstly, I created a new queue and I wanted to figure out how much speed I will get with for example PCQ discipline.

```
queue add name=dscp-queue target=10.0.30.2 dst=10.0.10.2 queue=pcq-upload-default/pcql-
```

The screenshot shows the Winbox interface on a MikroTik device. In the terminal window, the command `queue simple print where name=dscp-queue` is run, displaying the configuration for a DSCP queue named "dscp-queue". The queue is set to target IP 10.0.30.2/32 and destination IP 10.0.10.2/32, with priority 8/8. It uses the "pcq-upload-default/pcq-download-default" discipline, has a limit-at of 0/0, a max-limit of 1M/1M, a burst-limit of 0/0, a burst-threshold of 0/0, a burst-time of 0s/0s, and a bucket-size of 0.1/0.1. The Winbox interface also shows a network diagram with nodes PC1, PC3, Router, and Switch1.

```
[admin@MikroTik] > queue simple print where name=dscp-queue
Flags: X - disabled, I - invalid; D - dynamic
3  name="dscp-queue" target=10.0.30.2/32 dst=10.0.10.2/32 parent=none packet-marks="" priority=8/8
queue=pcq-upload-default/pcq-download-default limit-at=0/0 max-limit=1M/1M burst-limit=0/0 burst-threshold=0/0 burst-time=0s/0s bucket-size=0.1/0.1
[admin@MikroTik] >
```

- However, it gave me 0 speed when running together with the iperf3

The screenshot shows the Winbox interface with multiple terminal windows. One window runs `curl --local-port 42154-42155 -o ftp://ubuntu:ubuntu@10.0.10.2/my_file` to download a file from PC3. Another window shows the download progress with a table of statistics. The network diagram indicates traffic flow from PC1 through a switch to PC3. The Winbox interface also shows a tree view of network nodes including Cloud, PC1, PC2, PC3, PC4, Router, and various switches.

```
ubuntu@ubuntu-cloud:~$ curl --local-port 42154-42155 -o ftp://ubuntu:ubuntu@10.0.10.2/my_file
% Total    % Received % Xferd  Average Speed   Time     Time   Current
          Dload  Upload   Total Spent  Spent  Left  Speed
 1 50.0M   1 755k   0  0:07:16  0:00:06  0:07:10  118k
 1 50.0M   1 999k   0  0:07:24  0:00:08  0:07:16  114k
 3 50.0M   3 1737k   0  0:07:56  0:00:16  0:07:40  90342
 3 50.0M   3 1737k   0  0:08:26  0:00:17  0:08:09  68495
 3 50.0M   3 1737k   0  0:08:57  0:00:18  0:08:39  47700
 3 50.0M   3 1737k   0  0:09:27  0:00:19  0:09:08  26616
 3 50.0M   3 1737k   0  0:09:57  0:00:20  0:09:37  5220
 3 50.0M   3 1737k   0  0:10:27  0:00:21  0:10:06  0
 3 50.0M   3 1737k   0  0:10:57  0:00:22  0:10:35  0
 3 50.0M   3 1737k   0  0:14:28  0:00:29  0:13:59  0 10.0.30.2/24
 4 50.0M   4 2313k   0  0:21:57  0:00:59  0:20:58  122k
 4 50.0M   4 2420k   0  0:21:18  0:01:00  0:20:18  118k
 5 50.0M   5 2658k   0  0:20:02  0:01:02  0:19:00  118k^C
ubuntu@ubuntu-cloud:~$
```

- Therefore I created a new queue type based on CAKE queueing discipline to use it with DSCP

```
queue type add name=my-cake-queue kind cake cake-diffserv=diffserv4
queue simple set dscp-queue queue=my-cake-queue/my-cake-queue
```

The screenshot shows the Winbox interface with a terminal window running `queue type print where name=my-cake-queue`. It displays the configuration for a CAKE queue named "my-cake-queue" with kind=cake, bandwidth=0bps, overhead=0, overhead-scheme="", RTT=100ms, and diffserv4. Another window shows the configuration for a DSCP queue named "dscp-queue" with priority 8/8, target 10.0.30.2/32, and destination 10.0.10.2/32, using the CAKE queue as its parent. The Winbox interface also shows a network diagram and a tree view of nodes.

```
[admin@MikroTik] > queue type print where name=my-cake-queue
Flags: * - default
5  name="my-cake-queue" kind=cake cake-bandwidth=0bps cake-overhead=0 cake-overhead-scheme="" cake-rtt=100ms cake-diffserv=diffserv4
cake-flowmode=triple-isolate cake-nat=no cake-wash=no cake-ack-filter=none
[admin@MikroTik] > queue simple print where name=dscp-queue
Flags: X - disabled, I - invalid; D - dynamic
3  name="dscp-queue" target=10.0.30.2/32 dst=10.0.10.2/32 parent=none packet-marks="" priority=8/8 queue=my-cake-queue/my-cake-queue
limit-at=0/0 max-limit=1M/1M burst-limit=0/0 burst-threshold=0/0 burst-time=0s/0s bucket-size=0.1/0.1
[admin@MikroTik] >
```

- And of course I needed to configure marking FTP packets with DSCP values, so I configured them to have DSCP = 63 which is the maximum possible DSCP value

```
ip firewall mangle add chain=prerouting action=change-dscp new-dscp=63 protocol=tcp src-address=10.0.30.2 src-port=42154
ip firewall mangle add chain=prerouting action=change-dscp new-dscp=63 protocol=tcp src-address=10.0.30.2 src-port=42155
ip firewall mangle add chain=prerouting action=change-dscp new-dscp=63 protocol=tcp dst-address=10.0.30.2 dst-port=42154
ip firewall mangle add chain=prerouting action=change-dscp new-dscp=63 protocol=tcp dst-address=10.0.30.2 dst-port=42155
```

```

[admin@MikroTik] > ip firewall mangle print
Flags: X - disabled, I - invalid; D - dynamic
0 chain=prerouting action=mark-packet new-packet-mark=ftp-alloc protocol=tcp src-port=42154
1 chain=prerouting action=mark-packet new-packet-mark=ftp-alloc protocol=tcp src-port=42155
2 chain=prerouting action=mark-packet new-packet-mark=ftp-alloc protocol=tcp dst-port=42154
3 chain=prerouting action=mark-packet new-packet-mark=ftp-alloc protocol=tcp dst-port=42155
4 chain=prerouting action=change-dscp new-dscp=63 protocol=tcp src-address=10.0.30.2 src-port=42154
5 chain=prerouting action=change-dscp new-dscp=63 protocol=tcp src-address=10.0.30.2 src-port=42155
6 chain=prerouting action=change-dscp new-dscp=63 protocol=tcp dst-address=10.0.30.2 dst-port=42155
7 chain=prerouting action=change-dscp new-dscp=63 protocol=tcp dst-address=10.0.30.2 dst-port=42154
[admin@MikroTik] >

```

- As a result the speed did reach 0, so it seems that it worked!

```

ubuntu@ubuntu-cloud:~$ curl --local-port 42154-42155 -O ftp://ubuntu:ubuntu@10.0.10.2/my_file
% Total    % Received % Xferd  Average Speed   Time     Time      Current
          Dload  Upload Total Spent   Spent    Left Speed
0 50.0M    0 449k    0  0:07:14  0:00:03  0:07:11  117k
1 50.0M    1 564k    0  0:07:14  0:00:04  0:07:10  117k
2 50.0M    2 1111k   0  0:08:17  0:00:10  0:08:07  87546
2 50.0M    2 1165k   0  0:08:38  0:00:11  0:08:27  74056
2 50.0M    2 1224k   0  0:08:54  0:00:12  0:08:42  61999
2 50.0M    2 1283k   0  0:09:26  0:00:13  0:08:57  59463
2 50.0M    2 1404k   0  0:09:38  0:00:15  0:09:23  58910
3 50.0M    3 1633k   0  0:10:22  0:00:19  0:10:03  59125
3 50.0M    3 1802k   0  0:10:49  0:00:22  0:10:27  58331
3 50.0M    3 1916k   0  0:11:02  0:00:24  0:10:38  58398 10.0.30.0/24
4 50.0M    4 2150k   0  0:11:26  0:00:28  0:10:58  59320
5 50.0M    5 2888k   0  0:11:45  0:00:39  0:11:06  79531 switch3
5 50.0M    5 3009k   0  0:11:34  0:00:40  0:10:54  92450
6 50.0M    6 3126k   0  0:11:24  0:00:41  0:10:43  102k
6 50.0M    6 3245k   0  0:11:15  0:00:42  0:10:33  114k
6 50.0M    6 3362k   0  0:11:06  0:00:43  0:10:23  118k
7 50.0M    7 3601k  ~ 0  0:10:59  0:00:44  0:10:15  118k
ubuntu@ubuntu-cloud:~$

```

7. Choice and install any tool that you like for bandwidth control/netflow analysis/network control & monitoring. Play around with the network settings and show the different QoS metrics via UI

8. Try to answer the question: packet drops can occur even in an unloaded network where there is no queue overflow. In what cases and why does this happen?

There are several cases where it is possible that I can figure out:

- Firewalls** can apply special filtering for the packets and it might be not obvious for the sender what's happening.
- There can be **errors on physical layer** that can lead to packets corruption. This may look as a packet loss from both sides.
- TTL expiration**. Packets have TTL which is measured as number of hops usually. When TTL reaches zero, a packet is dropped and appropriate ICMP message is sent to the sender.
- RED scheduling algorithm** can be applied for queuing or a similar algorithm. It has a certain probability to drop packets even if the queue is not full. This is done to prevent the overflow.

Task 3 - QoS verification & packets analysis

1. How can you check if your QoS rules are applied correctly? List and describe the various methods.

I am not sure of what is required, so I just can suggest such methods:

- Traffic Generators and Test Tools**. Utilizing traffic generation tools (such as `iperf` or `Ostinato`) to simulate different types of traffic on your network. This can help you understand how QoS rules affect performance in a

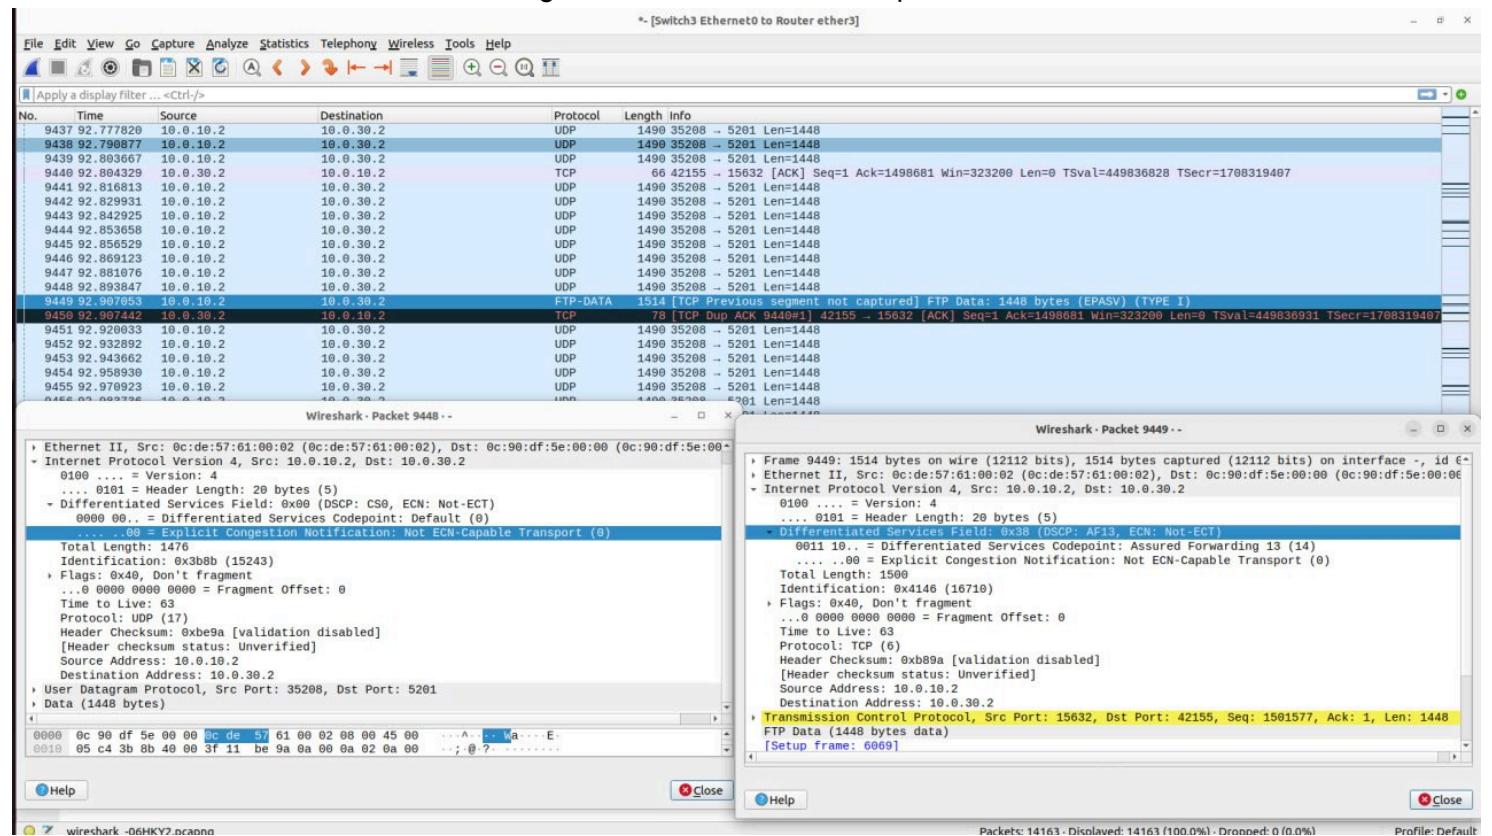
testing environment. Of course, it is better not to use it in production environment together with the production load. The method is pretty simple: generate traffic and verify its distribution with the expected values.

- **Network flow analysis.** We can monitor on how traffic is distributed in production environment by exploring it using different analytical tools, network bitrate graphs and etc, e.g. `iperf`.
- **Traffic Monitoring Tools.** Use network monitoring tools to observe real-time traffic flow and check whether the QoS rules are being enforced on the network. Tools like **Wireshark** can capture and help us analysing traffic. For example, in filtering traffic based on the QoS parameters (like DSCP values) to check if the packets are marked and prioritised correctly.
- **Network Device Interfaces.** Check the configuration and status of network devices (routers, switches) where QoS is configured. For example, we can monitor if packets marking is happening using device's monitoring tools or accumulated statistics. For example, in Mikrotik it is possible to look for statistics on dropped packets, queue lengths, and traffic distribution.

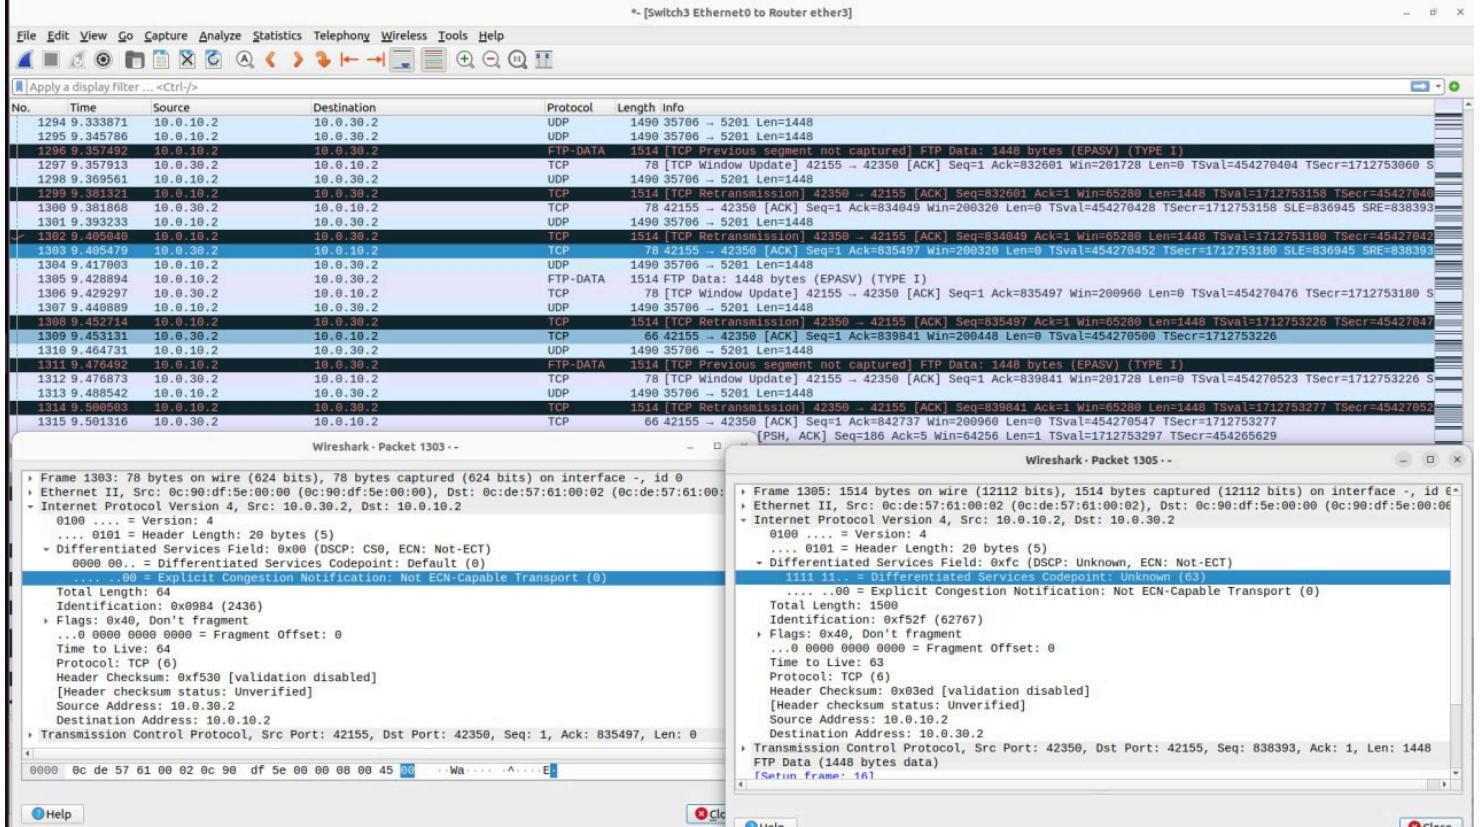
2. Try to use Wireshark to see the QoS packets. How does this depend on the number of routers in the network topology?

It really depends on the method of QoS used. For example, we wouldn't find anything interesting in Wireshark if traffic allocation technique is used. However, for DSCP case we can find.

- Let me show DSCP values that were assigned to FTP traffic from the previous task.



- Here on the left we see a UDP datagram that was generated by `iperf3` and it had the default DSCP value = 0. However, on the right hand-side we can notice our FTP packet which is marked with value = 63



References

- <https://help.mikrotik.com/docs/spaces/ROS/pages/18350234/Cloud+Hosted+Router+CHR>
- <https://help.mikrotik.com/docs/spaces/ROS/pages/7962644/Bandwidth+Test>
- <https://d2cpnw0u24fjm4.cloudfront.net/wp-content/uploads/iPerf3-User-Documentation.pdf>